

# **A Monte-Carlo Simulator For Light Propagation in Turbid Media using HLS**

**Abdul-Amir Yassine**

**Omar Ismail**

**Yasmin Afsharnejad**



# Outline

- Motivation
- Algorithm Overview
- Design Overview
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- Verification and Testing Procedure
- Results and Resource Utilization
- Experience with HLS

# Outline


- **Motivation**
- Algorithm Overview
- Design Overview
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- Verification and Testing Procedure
- Results and Resource Utilization
- Experience with HLS

# Motivation

- Emerging clinical applications require **fast** and **accurate** modelling of **light propagation** in turbid media.
  - Bioluminescence Imaging (BLI).
  - Photodynamic Therapy (PDT).
- Radiative transfer equation (RTE) is **computationally expensive**.
- Monte Carlo simulations can model the problem **accurately**, but they are **slow**!
- **High-level synthesis** facilitated the acceleration of such methods in recent years.

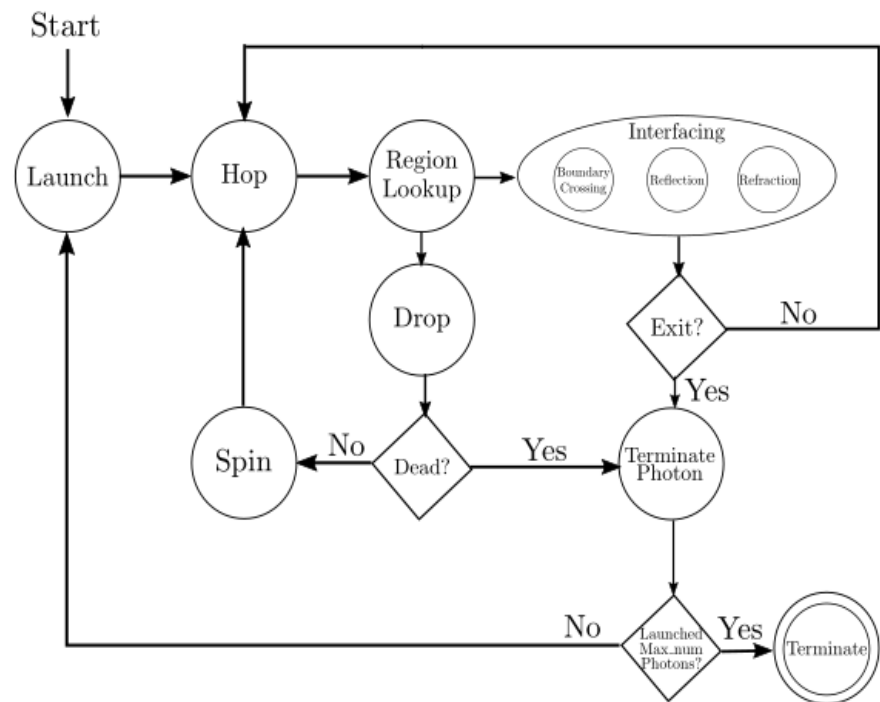
# Problem Definition

- Given:
  - Geometry specification (3D tissues with a voxelized representation).
  - Material optical properties.
  - Light source position. (Point Light Sources)
- Output:
  - Light absorption in voxel elements.
  - Trace of light in the tissue.



# Algorithm Stages

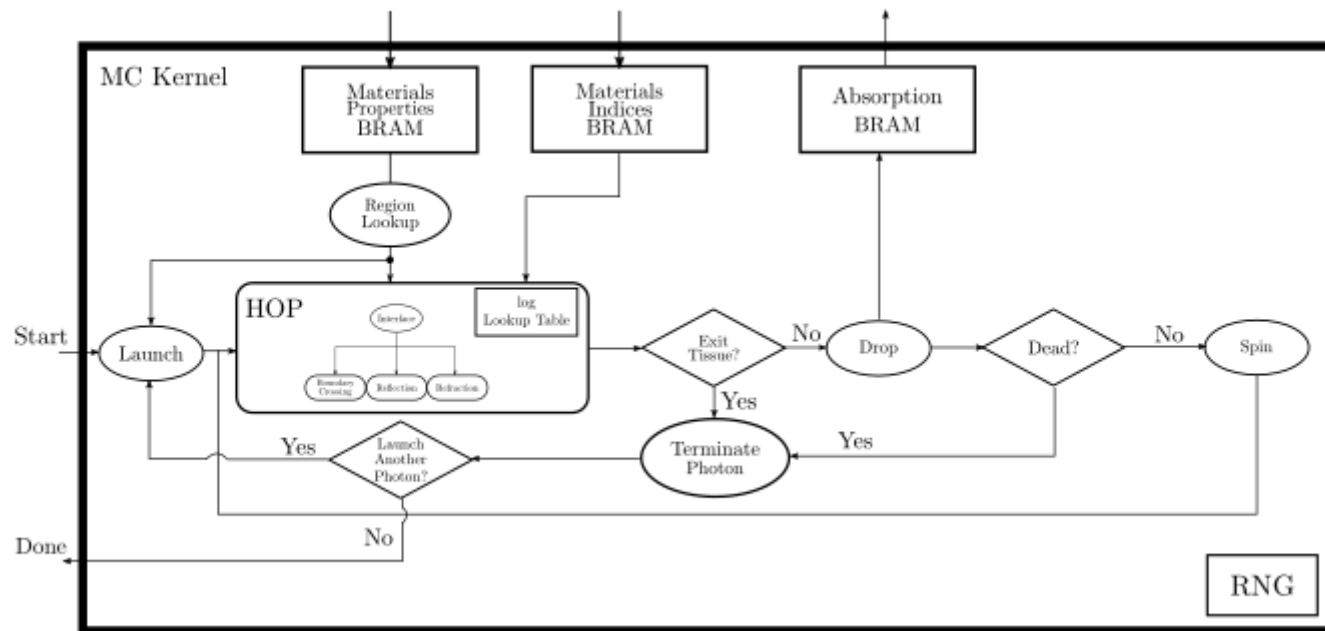
- **Launch:** Generate a packet of photons with a random direction from the source position.
- **Hop:** Moves the simulated photon from current position to a new position.
- **Drop:** Computes the photon's weight to **simulate the energy absorbed** by the tissue based on the absorption coefficient.
- **Roulette:** Checks if the photon packet is dead.
- **Spin:** Deriving the **scattering angle** for the next hop.



# Outline

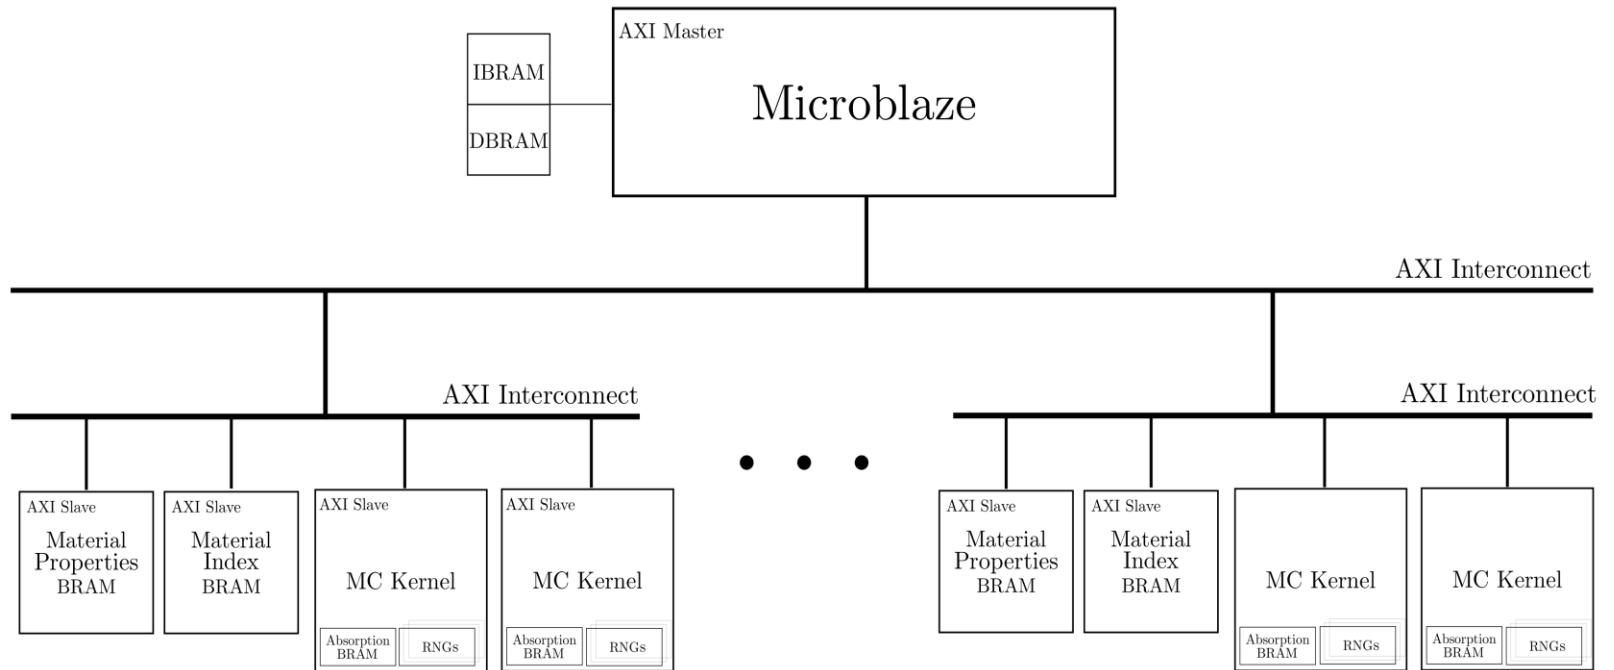
- Motivation
- Algorithm Overview
- **Design Overview**
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- Verification and Testing Procedure
- Results and Resource Utilization
- Experience with HLS

# FSM







# Final Design



# Big Picture!



# RNG

- TinyMT:
  - A smaller implementation of Mersenne Twister<sup>[1]</sup>.
  - State is 127bits long → Period is  $2^{127} - 1$ .
  - Only bitwise operations → easily implemented on FPGAs and fast.

```
void tinymt32_next_state()
{
#pragma HLS INLINE

    unsigned int x;
    unsigned int y;

    y = status[3];
    x = (status[0] & (unsigned int)MY_TINYMT32_MASK)
    ^ status[1]
    ^ status[2];
    x ^= (x << (unsigned int)MY_TINYMT32_SH0);
    y ^= (y >> (unsigned int)MY_TINYMT32_SH0) ^ x;

    status[0] = status[1];
    status[1] = status[2] ^ (-((int)(y & 1)) & mat1);
    status[2] = x ^ (y << (unsigned int)MY_TINYMT32_SH1)
    ^ (-((int)(y & 1)) & mat2);
    status[3] = y;
}
```

[1] Matsumoto, M. and Nishimura, T., 1998. Mersenne twister: a 623-dimensionally equidistributed uniform pseudo-random number generator. ACM Transactions on Modeling and Computer Simulation (TOMACS), 8(1), pp.3-30.

# Outline

- Motivation
- Algorithm Overview
- Design Overview
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- Verification and Testing Procedure
- Results and Resource Utilization
- Experience with HLS

# Pipelining Challenges

- Each photon packet runs **independently** from the other.
- Each photon packet depends on its **previous** iteration.
- Ideally, datapath can be **pipelined** . 😊
- **Problems:** 😞
  - Boundary crossing.
  - RNG dependency.
  - Absorption Accumulation.

# Pipelining Challenges (Cont'd)

- Boundary Crossing:
  - Every time a photon crosses a boundary, it goes through **HOP again** to jump remaining step.
  - **Problem:** Pipeline has to **stall** as following photons **cannot** enter the Drop stage **until the current photon** does.
  - **Solution:** We **removed HOP while loop** and prevented photon from going to the Drop stage unless a maximum number of crossings, or no crossing occurs.

```
PHOTON_LOOP1: for ( int i = 0; i < NUM_PACKETS_LAUNCH; i++)
{
#pragma HLS PIPELINE II=8

    ph = sim1->launch(i);

    bool res = sim1->HOP( i );
    if (res)
    {
        sim1->terminate();
        sim1->photons[i] = ph;
        iter_idx1++;
    }
    if (sim1->get_photon_i_num_hits(i) >= MAX_NUM_HITS_HOP)
    {
        sim1->set_photon_i_remaining_step(i , 0);
        sim1->SPIN(i);
        if (sim1->DROP(i)){
            sim1->photons[i] = ph;
            iter_idx1++;
        }
        sim1->set_photon_i_num_of_hits(i, 0);
    }
}
```

# Pipelining Challenges (Cont'd)

- Random Number Generator Dependency:
  - Every stage needs an RNG.
  - **Problem: Different photons** in different stages of the pipeline **access** the RNG.
  - **Solution:** Assign every stage a **different RNG** with different seed.
- Absorption Accumulation:
  - Photon's partial energy weight should be deposited in tissue.
  - **Problem: RAW dependence:**
    - Need to read the current value, accumulate then store before other photons can do the same.
    - Takes 6 cycles.
  - **Solution:** Partial accumulation.
    - Accumulate in different buffers then add at the end.
  - **Problem:** Large BRAM utilization due to large data sizes.

# Float Tricks

```
typedef union {  
    float_used f;  
    int ui;  
} ftoi_union;
```

- Multiplication/Division by 2:

```
// generating random initial direction  
ftoi_union rand_z_1, rand_pi_2;  
rand_z_1.f = _random_generator.tinynt32_generate_float();  
rand_pi_2.f = PI*_random_generator.tinynt32_generate_float();  
  
// Multiplying the random number by 2  
rand_z_1.ui += NUM_TO_ADD_UNION;  
rand_pi_2.ui += NUM_TO_ADD_UNION;
```

- Subtraction of 1 from a number between 1 and 2.
- Float Comparison.
- Division by constant: store inverse at the beginning and multiply.



# Other Optimization

- **Log** was on critical path.
  - Input values are only between 0 and 1.
  - **Approximate** by a **large lookup table** that is distributed uniformly.
- **HLS didn't detect** similar or independent code in different branches.
  - **Move code** outside of branches to reduce branch divergence.
  - Example:


```
sim1->set_photon_i_remaining_step(i , 0);  
sim1->SPIN(i);  
if (sim1->DROP(i)){  
    sim1->photons[i] = ph;  
    iter_idx1++;  
}  
sim1->set_photon_i_num_of_hits(i, 0);
```

# Final Result

- All features are supported.
  - Boundary Crossing.
  - Reflection (internal and Fresnel's).
  - Refraction.
  - Absorption.
- Optimal **II is 6**.
- With II of 6, timing met was 10ns.
- By changing target **II to 8**, we were able to meet **6ns** (initial target was 5ns).
- This translates to **167MHz**.

# Challenges in Integration

- The eight-core design consists of **large number** of slaves and masters.
- Problems:
  - **AXI\_Interconnects** have limits in number of slaves and masters.
  - **Timings** were **not met** in eight core implementation.
- Solution:
  - **Cascading** the AXI-Interconnect.



- **Adding register** slices on interconnect CP to meet the target frequency.

# Outline

- Motivation
- Algorithm Overview
- Design Overview
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- **Verification and Testing Procedure**
- Results and Resource Utilization
- Experience with HLS

# Testing


- **Simulation Test-bench in Vivado HLS** for getting the absorption results.
  - We wrote **scripts** to:
    - Generate random data with different material properties.
    - Convert results to VTK format to visualize them.
  - We compared results to **existing tools** to check output consistency.
- Verification was done through **RTL simulation**.
- After Integration in Vivado, we verified the results using **behavioral Simulation waveforms**.
- After moving the design to board, **SDK system debugger** was used to dump results to files.

# Outline

- Motivation
- Algorithm Overview
- Design Overview
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- Verification and Testing Procedure
- **Results and Resource Utilization**
- Experience with HLS

# Results

- Test case:
  - FPGA: Xilinx® Kintex® Ultrascale™ XCKU115-2 - FLVA1517E (20 nm).
  - One million photons simulated.
  - Tissue size is 8K elements.
  - Eight Kernels.
- CPU time: 14m 30 seconds
- Hardware time: 5 seconds
- Speed-up: 174x
- FullMonte's Time: 16 secs



# Resource Utilization

Resource	Available	Utilization	Utilization (%)
LUTs	663K	299K	45
LUTRAMs	293K	20K	7
Registers	1.326M	328K	25
BRAMs	2160	1206	55
DSPs	5520	731	13



# Outline

- Motivation
- Algorithm Overview
- Design Overview
- Implementation Challenges
  - Pipelining
  - Float Tricks and Other Optimizations
  - Integration
- Verification and Testing Procedure
- Results and Resource Utilization
- Experience with HLS

# Experience with HLS

- **Simple**, easy to use.
- No need for low level hardware experience.
- It would have taken us at least **SIX Months in HDL!**
- Issues:
  - **HLS hates if-statements!!**
  - Couldn't schedule some independent sections of the code in parallel.
    - Example: handling multiple kernels inside HLS.
  - Minimal support for floating point optimizations.

Thank You!